STAT 401A - Statistical Methods for Research Workers
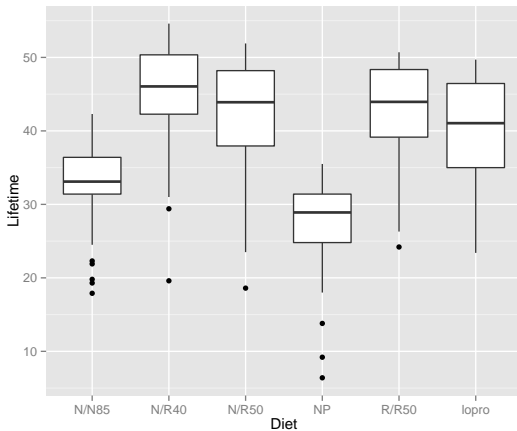## One-way ANOVA (contrasts and multiple comparisons)

Jarad Niemi (Dr. J)

Iowa State University

last updated: October 9, 2014

Mice lifetimes

Simple hypothesis

Consider the one-way ANOVA model: $Y_{ij} \sim N(\mu_j, \sigma^2)$ where $j = 1, \ldots, J$.

Here are a few simple alternative hypotheses:

1. Mean lifetimes for N/R50 and R/R50 diet are different.
2. Mean lifetimes for N/R40 is different than for N/R50 and R/R50 combined.
3. Mean lifetimes for high calorie (NP and N/N85) diets is different than for low calorie diets combined.

$H_0 : \gamma = 0 \qquad H_1 : \gamma \neq 0 :$

$$
\begin{aligned}
\gamma_1 &= \mu_{R/R50} - \mu_{N/R50} \\
\gamma_2 &= \mu_{N/R40} - \tfrac{1}{2}(\mu_{N/R50} + \mu_{R/R50}) \\
\gamma_3 &= \tfrac{1}{4}(\mu_{N/R50} + \mu_{R/R50} + \mu_{N/R40} + \mu_{lopro}) - \tfrac{1}{2}(\mu_{NP} + \mu_{N/N85})
\end{aligned}
$$

Contrasts

Definition

A linear combination of group means has the form

$$\gamma = C_1\mu_1 + C_2\mu_2 + \ldots + C_J\mu_J$$

where C_j are known coefficients and μ_j are the unknown population means.

Definition

A linear combination with $C_1 + C_2 + \cdots + C_J = 0$ is a contrast.

Remark Contrast interpretation is usually best if
$|C_1| + |C_2| + \cdots + |C_J| = 2$, i.e. the positive sum to 1 and the negative coefficients sum to -1.

Inference on contrasts

$$\gamma = C_1\mu_1 + C_2\mu_2 + \cdots + C_J\mu_J$$

Estimated by

$$g = C_1\overline{Y}_1 + C_2\overline{Y}_2 + \cdots + C_J\overline{Y}_J$$

with standard error

$$SE(g) = s_p\sqrt{\frac{C_1^2}{n_1} + \frac{C_2^2}{n_2} + \cdots + \frac{C_J^2}{n_J}}$$

t-statistic (compare to t_{n-J}) and CI:

$$t = \frac{g}{SE(g)} \qquad g \pm t_{n-J}(1 - \alpha/2)SE(g)$$

Contrasts for mice lifetime dataset

For these contrasts:

1. Mean lifetimes for N/R50 and R/R50 diet are different.
2. Mean lifetimes for N/R40 is different than for N/R50 and R/R50 combined.
3. Mean lifetimes for high calorie (NP and N/N85) diets is different than for low calorie diets combined.

$H_0 : \gamma = 0$ $H_1 : \gamma \neq 0$:

$$
\begin{aligned}
\gamma_1 &= \mu_{R/R50} - \mu_{N/R50} \\
\gamma_2 &= \mu_{N/R40} - \tfrac{1}{2}(\mu_{N/R50} + \mu_{R/R50}) \\
\gamma_3 &= \tfrac{1}{4}(\mu_{N/R50} + \mu_{R/R50} + \mu_{N/R40} + \mu_{lopro}) - \tfrac{1}{2}(\mu_{NP} + \mu_{N/N85})
\end{aligned}
$$

	N/N85	N/R40	N/R50	NP	R/R50	lopro
early rest - none @ 50kcal	0.00	0.00	-1.00	0.00	1.00	0.00
40kcal/week - 50kcal/week	0.00	1.00	-0.50	0.00	-0.50	0.00
lo cal - hi cal	-0.50	0.25	0.25	-0.50	0.25	0.25

Mice liftime examples

	Diet	n	mean	sd
1	N/N85	57	32.69	5.13
2	N/R40	60	45.12	6.70
3	N/R50	71	42.30	7.77
4	NP	49	27.40	6.13
5	R/R50	56	42.89	6.68
6	lopro	56	39.69	6.99

Contrasts:

	g	SE(g)	t	p	L	U
early rest - none @ 50kcal	0.59	1.19	0.49	0.62	-1.76	2.94
40kcal/week - 50kcal/week	2.53	1.05	2.41	0.02	0.46	4.59
lo cal - hi cal	12.45	0.78	15.96	0.00	10.92	13.98

```
DATA case0501;
   INFILE 'case0501.csv' DSD FIRSTOBS=2;
   INPUT lifetime diet $ ;

PROC MEANS DATA=case0501;
   CLASS diet;
   VAR lifetime;
   RUN;
```

The MEANS Procedure
Analysis Variable : lifetime

diet	N Obs	N	Mean	Std Dev	Minimum	Maximum
N/N85	57	57	32.6912281	5.1252972	17.9000000	42.3000000
N/R40	60	60	45.1166667	6.7034058	19.6000000	54.6000000
N/R50	71	71	42.2971831	7.7681947	18.6000000	51.9000000
NP	49	49	27.4020408	6.1337010	6.4000000	35.5000000
R/R50	56	56	42.8857143	6.6831519	24.2000000	50.7000000
lopro	56	56	39.6857143	6.9916945	23.4000000	49.7000000

```
PROC GLM;
  CLASS diet;
  MODEL lifetime = diet / CLPARM;
  ESTIMATE 'early rest - none @ 50kcal' diet  0  1 -1  0  0  0 ;
  ESTIMATE '40kcal/week - 50kcal/week'  diet  0  2 -1  0 -1  0 / DIVISOR = 2;
  ESTIMATE 'lo cal - hi cal'            diet -2  1  1 -2  1  1 / DIVISOR = 4 ;
  RUN;
  QUIT;
```

                         The GLM Procedure

                                 Sum of
    Source   DF   Squares   Mean Square   F Value   Pr > F
    Model   5   12733.94181   2546.78836   57.10   <.0001
    Error   343   15297.41532   44.59888
    Corrected Total   348   28031.35713

                                               Standard
     Parameter   Estimate   Error   t Value   Pr > |t|
     early rest - none @ 50kcal   0.5885312   1.19355007   0.49   0.6223
     40kcal/week - 50kcal/week   2.5252180   1.04854904   2.41   0.0166
     lo cal - hi cal   12.4496851   0.78001425   15.96   <.0001

                Parameter   95% Confidence Limits
                early rest - none @ 50kcal   -1.7590676   2.9361299
                40kcal/week - 50kcal/week   0.4628224   4.5876136
                lo cal - hi cal   10.9154718   13.9838985

```
library(multcomp)
m = lm(Lifetime~Diet-1, case0501) # The -1 indicates no intercept (see Ch 7)
summary(m)




Call:
lm(formula = Lifetime ~ Diet - 1, data = case0501)

Residuals:
    Min      1Q  Median      3Q     Max
-25.517  -3.386   0.814   5.183  10.014

Coefficients:
           Estimate Std. Error t value Pr(>|t|)
DietN/N85    32.691      0.885    37.0   <2e-16 ***
DietN/R40    45.117      0.862    52.3   <2e-16 ***
DietN/R50    42.297      0.793    53.4   <2e-16 ***
DietNP       27.402      0.954    28.7   <2e-16 ***
DietR/R50    42.886      0.892    48.1   <2e-16 ***
Dietlopro    39.686      0.892    44.5   <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.68 on 343 degrees of freedom
Multiple R-squared: 0.972,  Adjusted R-squared: 0.972
F-statistic: 2.01e+03 on 6 and 343 DF,  p-value: <2e-16


K


                          N/N85 N/R40 N/R50   NP R/R50 lopro
early rest - none @ 50kcal  0.0  0.00 -1.00  0.0  1.00  0.00
40kcal/week - 50kcal/week   0.0  1.00 -0.50  0.0 -0.50  0.00
lo cal - hi cal            -0.5  0.25  0.25 -0.5  0.25  0.25
```

```
t = glht(m, linfct=K)
summary(t)


 Simultaneous Tests for General Linear Hypotheses

Fit: lm(formula = Lifetime ~ Diet - 1, data = case0501)

Linear Hypotheses:
                             Estimate Std. Error t value Pr(>|t|)
early rest - none @ 50kcal == 0    0.589      1.194    0.49    0.946
40kcal/week - 50kcal/week == 0     2.525      1.049    2.41    0.049 *
lo cal - hi cal == 0              12.450      0.780   15.96   <1e-04 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
(Adjusted p values reported -- single-step method)


confint(t, calpha=univariate_calpha())


 Simultaneous Confidence Intervals

Fit: lm(formula = Lifetime ~ Diet - 1, data = case0501)

Quantile = 1.967
95% confidence level


Linear Hypotheses:
                             Estimate   lwr    upr
early rest - none @ 50kcal == 0  0.589  -1.759  2.936
40kcal/week - 50kcal/week == 0   2.525   0.463  4.588
lo cal - hi cal == 0            12.450  10.915 13.984
```

Summary

- Contrasts are linear combinations that sum to zero
- t-test tools are used to calculate pvalues and confidence intervals

SAS code and output for one-way ANOVA

```
DATA mice;
  INFILE 'case0501.csv' DSD FIRSTOBS=2;
  INPUT lifetime diet $;

PROC GLM DATA=mice;
  CLASS diet;
  MODEL lifetime = diet;
  LSMEANS diet / ADJUST=T;
  RUN;
```

The GLM Procedure

Dependent Variable: lifetime

Source	DF	Sum of Squares	Mean Square	F Value	Pr > F
Model	5	12733.94181	2546.78836	57.10	<.0001
Error	343	15297.41532	44.59888		
Corrected Total	348	28031.35713			

SAS code and output for pairwise comparisons

```
                      The GLM Procedure
                    Least Squares Means

                          lifetime        LSMEAN
                diet        LSMEAN         Number

                N/N85     32.6912281            1
                N/R40     45.1166667            2
                N/R50     42.2971831            3
                NP        27.4020408            4
                R/R50     42.8857143            5
                lopro     39.6857143            6


              Least Squares Means for effect diet
              Pr > |t| for H0: LSMean(i)=LSMean(j)

                  Dependent Variable: lifetime
```

i/j	1	2	3	4	5	6
1		<.0001	<.0001	<.0001	<.0001	<.0001
2	<.0001		0.0166	<.0001	0.0731	<.0001
3	<.0001	0.0166		<.0001	0.6223	0.0293
4	<.0001	<.0001	<.0001		<.0001	<.0001
5	<.0001	0.0731	0.6223	<.0001		0.0117
6	<.0001	<.0001	0.0293	<.0001	0.0117	

NOTE: To ensure overall protection level, only probabilities associated with pre-planned comparisons should be used.

Statistical testing errors

Definition
A type I error occurs when a true null hypothesis is rejected.

Definition
A type II error occurs when a false null hypothesis is not rejected. Power is one minus the type II error probability.

Remark We set α to control the type I error probability. If we set $\alpha = 0.05$, then we will incorrectly reject a true null hypothesis 5% of the time.

Definition
The familywise error rate is the probability of rejecting at least one true null hypothesis.

Type I error for all pairwise comparisons of J groups

How many combinations when choosing 2 items out of J?

$$\binom{J}{2} = \frac{J!}{2!(J-2)!}.$$

If $J = 6$, then there are 15 different comparison of means. If we set $\alpha = 0.05$ as our significance level, then individually each test will only incorrectly reject 5% of the time.

If we have 15 tests and use $\alpha = 0.05$, what is the familywise error rate?

$$1 - (1 - 0.05)^{15} = 1 - (0.95)^{15} = 1 - 0.46 = 0.54$$

So there is a greater than 50% probability of falsely rejecting a true null hypothesis!

Bonferroni correction

Definition

If we do m tests and want the familywise error rate to be α, the Bonferroni correction uses α/m for each individual test. The familywise error rate, for independent tests, is $1 - (1 - \alpha/, m)^m$.

Bonferroni familywise error rate
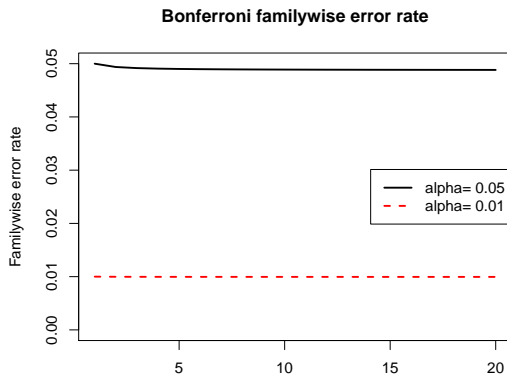
SAS code and output for pairwise comparisons

Compare the unadjusted pvalues to $\alpha/15 = 0.05/15 = 0.0033$.

```
                    Least Squares Means for effect diet
                    Pr > |t| for H0: LSMean(i)=LSMean(j)

                    Dependent Variable: lifetime

 i/j        1         2         3         4         5         6

  1                <.0001    <.0001    <.0001    <.0001    <.0001
  2      <.0001              0.0166    <.0001    0.0731    <.0001
  3      <.0001    0.0166              <.0001    0.6223    0.0293
  4      <.0001    <.0001    <.0001              <.0001    <.0001
  5      <.0001    0.0731    0.6223    <.0001              0.0117
  6      <.0001    <.0001    0.0293    <.0001    0.0117
```

NOTE: To ensure overall protection level, only probabilities associated with pre-planned
    comparisons should be used.

Now 2-3, 3-6, and 5-6 are no longer significant.

SAS code and output for one-way ANOVA

If you use SAS to do the adjustment, compare pvalues to $\alpha = 0.05$.

```
DATA mice;
  INFILE 'case0501.csv' DSD FIRSTOBS=2;
  INPUT lifetime diet $;

PROC GLM DATA=mice;
  CLASS diet;
  MODEL lifetime = diet;
  LSMEANS diet / ADJUST=BON;
  RUN;
```

```
                          Least Squares Means for effect diet
                          Pr > |t| for H0: LSMean(i)=LSMean(j)

                              Dependent Variable: lifetime

     i/j          1            2            3            4            5            6

      1                      <.0001       <.0001       0.0009       <.0001       <.0001
      2        <.0001                     0.2488       <.0001       1.0000       0.0002
      3        <.0001       0.2488                     <.0001       1.0000       0.4402
      4        0.0009       <.0001       <.0001                     <.0001       <.0001
      5        <.0001       1.0000       1.0000       <.0001                     0.1751
      6        <.0001       0.0002       0.4402       <.0001       0.1751
```

Comments on the Bonferroni correction

Remark The Bonferroni correction can be used in any situation. In particular, it can be used on unadjusted pvalues reported in an article that has many tests by comparing their pvalues to α/m where m is the number of tests they perform.

Remark The Bonferroni correction is (in general) the most conservative multiple comparison adjustment we will discuss, i.e. it will lead to the least null hypothesis rejections.

Constructing multiple confidence intervals

Remark A $100(1 - \alpha)$% confidence interval should contain the true value $100(1 - \alpha)$% of the time.

Remark An error occurs if the confidence interval does not contain the true value.

Just like the Type I error and familywise error rate, we can ask what is the probability at least one confidence interval does not cover the true value.

The procedures we will talk about for confidence intervals have equivalent approaches for hypothesis testing (pvalues). Within these procedures we still have the equivalence between pvalues and CIs.

Constructing multiple confidence intervals

Confidence interval for the difference between group j and group j':

$$\overline{Y}_j - \overline{Y}_{j'} \pm M \, s_p \sqrt{\frac{1}{n_j} + \frac{1}{n_{j'}}}$$

where M is a multiplier that depends on the adjustment procedure:

Procedure	M	Use
LSD	$t_{n-J}(1 - \alpha/2)$	After significant F-test (no adjustment)
Dunnett	multivariate t	Compare all groups to control
Tukey-Kramer	$q_{J,n-J}(1 - \alpha/2)/\sqrt{2}$	All pairwise comparisons
Scheffé	$\sqrt{(J-1)F_{(J-1,n-J)}(1 - \alpha)}$	All contrasts
Bonferroni	$t_{n-J}(1 - \alpha/2m)$	m tests (most generic)

SAS code and output for one-way ANOVA

```
DATA mice;
  INFILE 'case0501.csv' DSD FIRSTOBS=2;
  INPUT lifetime diet $;

PROC GLM DATA=mice;
  CLASS diet;
  MODEL lifetime = diet;
  LSMEANS diet / CL ADJUST=TUKEY;
  RUN;
                          The GLM Procedure
                        Least Squares Means
              Adjustment for Multiple Comparisons: Tukey-Kramer


                        Dependent Variable: lifetime

   i/j          1          2          3          4          5          6
    1                  <.0001     <.0001     0.0008     <.0001     <.0001
    2      <.0001                 0.1565     <.0001     0.4684     0.0002
    3      <.0001     0.1565                <.0001     0.9964     0.2460
    4      0.0008     <.0001     <.0001                <.0001     <.0001
    5      <.0001     0.4684     0.9964     <.0001                0.1168
    6      <.0001     0.0002     0.2460     <.0001     0.1168
```

SAS code and output for one-way ANOVA

```
                    The GLM Procedure
                 Least Squares Means
      Adjustment for Multiple Comparisons: Tukey-Kramer


             Least Squares Means for Effect diet

               Difference          Simultaneous 95%
               Between          Confidence Limits for
     i    j     Means           LSMean(i)-LSMean(j)
     1    2   -12.425439     -15.965442     -8.885435
     1    3    -9.605955     -13.009741     -6.202169
     1    4     5.289187       1.560626      9.017749
     1    5   -10.194486     -13.795557     -6.593416
     1    6    -6.994486     -10.595557     -3.393416
     2    3     2.819484      -0.536769      6.175736
     2    4    17.714626      14.029406     21.399846
     2    5     2.230952      -1.325223      5.787128
     2    6     5.430952       1.874777      8.987128
     3    4    14.895142      11.340571     18.449714
     3    5    -0.588531      -4.009133      2.832070
     3    6     2.611469      -0.809133      6.032070
     4    5   -15.483673     -19.227592    -11.739755
     4    6   -12.283673     -16.027592     -8.539755
     5    6     3.200000      -0.416969      6.816969
```

False Discovery Rate

Remark Not wanting to make a single mistake is pretty conservative.

In high-throughput fiels a more common multiple comparison adjustment is false discovery rate.

Definition
False discovery rate procedures try to control the expected proportion of incorrectly rejected null hypotheses.

How to incorporate multiple comparison adjustments

1. Determine what tests are going to be run (before looking at the data) or what confidence intervals are going to be constructed.
2. Determine which multiple comparison adjustment is the most relevant
3. Use/state that adjustment and interpret your results

Sulfur effect on scab disease in potatoes

The experiment was conducted to investigate the effect of sulfur on controlling scab disease in potatoes. There were seven treatments: control, plus spring and fall application of 300, 600, 1200 lbs/acre of sulfur. The response variable was percentage of the potato surface area covered with scab averaged over 100 random selected potatoes. A completely randomized design was used with 8 replications of the control and 4 replications of the other treatments.

Cochran and Cox. (1957) Experimental Design (2nd ed). pg96 and Agron. J. 80:712-718 (1988)

Scientific question:

- Does sulfur have any impact at all?
- Is there a difference between spring and fall?
- Is there an effect of increased sulfur (expect more sulfur causes less scab)?

Data

```
   inf trt row col
1    9  F3   4   1
2   12   O   4   2
3   18  S6   4   3
4   10 F12   4   4
5   24  S6   4   5
6   17 S12   4   6
7   30  S3   4   7
8   16  F6   4   8
9   10   O   3   1
10   7  S3   3   2
11   4 F12   3   3
12  10  F6   3   4
13  21  S3   3   5
14  24   O   3   6
15  29   O   3   7
16  12  S6   3   8
17   9  F3   2   1
18   7 S12   2   2
19  18  F6   2   3
20  30   O   2   4
21  18  F6   2   5
22  16 S12   2   6
23  16  F3   2   7
24   4 F12   2   8
25   9  S3   1   1
26  18   O   1   2
27  17 S12   1   3
28  19  S6   1   4
29  32   O   1   5
30   5 F12   1   6
```
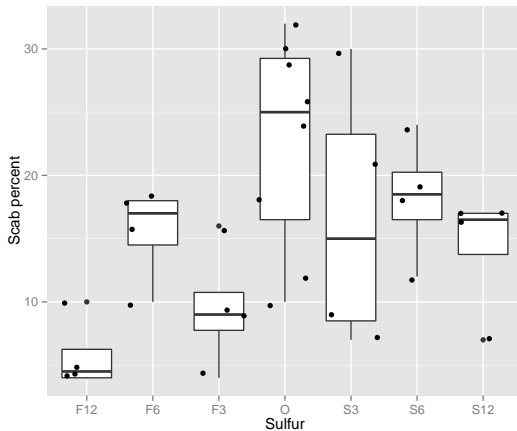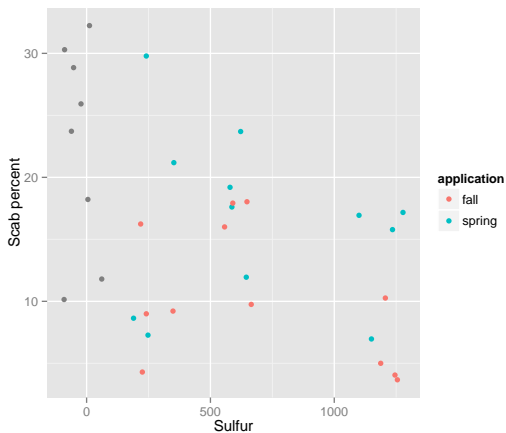
Design

Treatment visualization
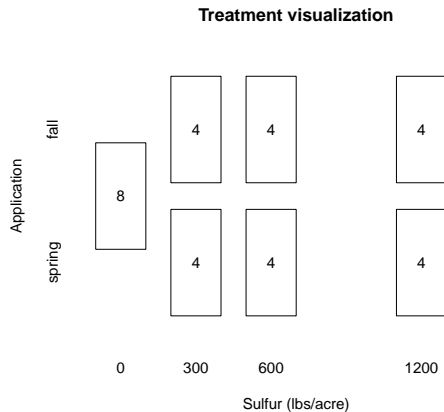
Design

**Completely randomized design
potato scab experiment**

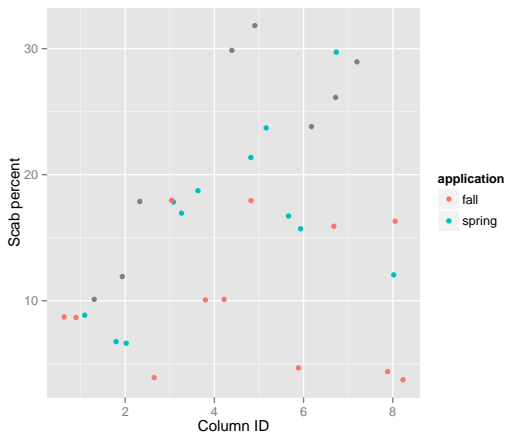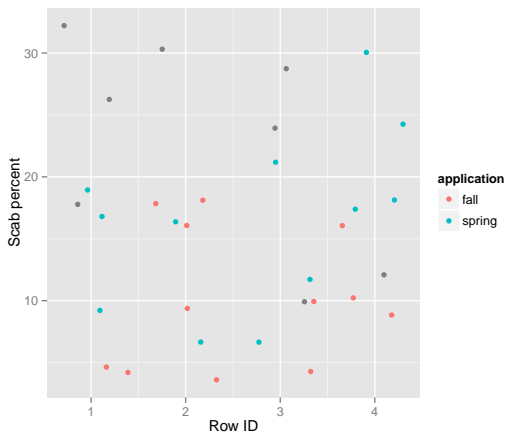|	1	2	3	4	5	6	7	8
4	F3	O	S6	F12	S6	S12	S3	F6
3	O	S3	F12	F6	S3	O	O	S6
2	F3	S12	F6	O	F6	S12	F3	F12
1	S3	O	S12	S6	O	F12	O	F3

row

col

Data

Data

Data

Data

Model

Y_{ij}: avg % of scab for observation i in treatment j for $j = 1, \ldots, 7$.

Assume $Y_{ij} \overset{ind}{\sim} N(\mu_j, \sigma^2)$.

Hypotheses:

- *Any effect*: One-way ANOVA F-test
- *Fall vs spring*: Contrast comparing fall vs spring applications
- *Sulfur level*: Linear trend contrast

Fall vs spring contrast

i	1	2	3	4	5	6	7	Div
Trt	F12	F6	F3	O	S3	S6	S12	
Fall v Spring	1	1	1	0	-1	-1	-1	3

- *Fall vs spring*: Contrast comparing fall vs spring applications

$$\gamma \quad = \frac{1}{3}(\mu_{F12} + \mu_{F6} + \mu_{F3}) + 0\mu_O - \frac{1}{3}(\mu_{S3} + \mu_{S6} + \mu_{S12})$$

$$= \frac{1}{3}\mu_{F12} + \frac{1}{3}\mu_{F6} + \frac{1}{3}\mu_{F3} + 0\mu_O - \frac{1}{3}\mu_{S3} - \frac{1}{3}\mu_{S6} - \frac{1}{3}\mu_{S12}$$

$$= \frac{1}{3}\left[\mu_{F12} + \mu_{F6} + \mu_{F3} + 0\mu_O - 1\mu_{S3} - 1\mu_{S6} - 1\mu_{S12}\right]$$

Linear trend contrasts

i	1	2	3	4	5	6	7	Div
Trt	F12	F6	F3	O	S3	S6	S12	
Fall v Spring	1	1	1	0	-1	-1	-1	3
Linear Trend	27	3	-9	-42	-9	3	27	8

- *Sulfur level*: Linear trend contrast
  - The unique X_i are 0, 3, 6, and 12.
  - So the linear trend contrast $(X_i - \overline{X})$ is
  - But 3, 6, and 12 are duplicated, so need the average of the groups

$$\gamma = -\frac{21}{4}\mu_0 - \frac{9}{4}\mu_3 + \frac{3}{4}\mu_6 + \frac{27}{4}\mu_{12}$$

$$= -\frac{21}{4}\mu_0 - \frac{9}{4}\left(\frac{\mu_{S3}+\mu_{F3}}{2}\right) + \frac{3}{4}\left(\frac{\mu_{S6}+\mu_{F6}}{2}\right) + \frac{27}{4}\left(\frac{\mu_{S12}+\mu_{F12}}{2}\right)$$

$$= \frac{1}{8}\left[-42\mu_0 - 9\mu_{S3} - 9\mu_{F3} + 3\mu_{S6} + 3\mu_{F6} + 27\mu_{S12} + 27\mu_{F12}\right]$$

SAS code

```
DATA d;
  INFILE 'potato.csv' DSD FIRSTOBS=2;
  INPUT scabp treatment $ row col;
  sulfur = 0;
  IF treatment in ("F3","S3") THEN sulfur=300;
  IF treatment in ("F6","S6") THEN sulfur=600;
  IF treatment in ("F12","S12") THEN sulfur=1200;
  application = "NA    ";
  IF treatment in ("F3","F6","F12") THEN application="fall";
  IF treatment in ("S3","S6","S12") THEN application="spring";

PROC PRINT DATA=d (OBS=10); RUN;

PROC MEANS;
  CLASS treatment;
  VAR scabp;
  RUN;
```

SAS code

```
     Obs    scabp    treatment    row    col    sulfur    application

      1        9         F3         4      1      300        fall
      2       12         O          4      2        0        NA
      3       18         S6         4      3      600        spring
      4       10         F12        4      4     1200        fall
      5       24         S6         4      5      600        spring
      6       17         S12        4      6     1200        spring
      7       30         S3         4      7      300        spring
      8       16         F6         4      8      600        fall
      9       10         O          3      1        0        NA
     10        7         S3         3      2      300        spring
```

```
                        The MEANS Procedure

                    Analysis Variable : scabp
```

treatment	N Obs	N	Mean	Std Dev	Minimum	Maximum
F12	4	4	5.7500000	2.8722813	4.0000000	10.0000000
F3	4	4	9.5000000	4.9328829	4.0000000	16.0000000
F6	4	4	15.5000000	3.7859389	10.0000000	18.0000000
O	8	8	22.6250000	8.3655330	10.0000000	32.0000000
S12	4	4	14.2500000	4.8562674	7.0000000	17.0000000
S3	4	4	16.7500000	10.7819293	7.0000000	30.0000000
S6	4	4	18.2500000	4.9244289	12.0000000	24.0000000

SAS code

```
PROC GLM DATA=d PLOTS=(DIAGNOSTICS RESIDUALS);
  CLASS treatment;
  MODEL scabp = treatment;
  LSMEANS treatment / CL;
  ESTIMATE 'fall - spring' treatment 1 1 1 0 -1 -1 -1 / DIVISOR=3;
  ESTIMATE 'linear trend'  treatment 27 -9 3 -42 27 -9 3 / DIVISOR=8;
  OUTPUT OUT=dres P=predicted R=residuals;
  RUN;

PROC GPLOT DATA=dres;
  PLOT residuals*predicted;
  PLOT residuals*sulfur;
  PLOT residuals*application;
  PLOT residuals*row;
  PLOT residuals*col;
  RUN;
```

Diagnostics
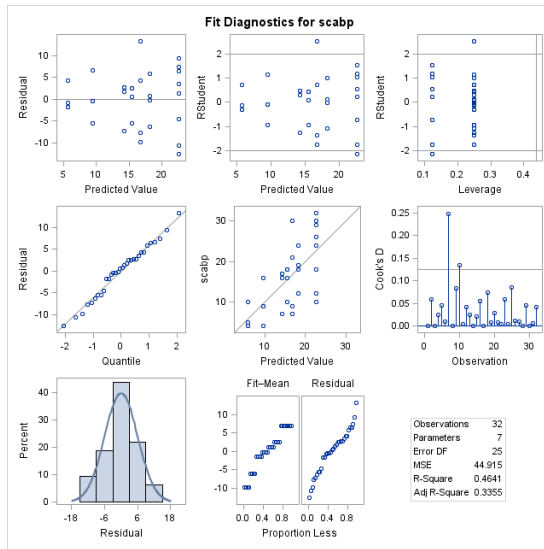
SAS output

```
                        The GLM Procedure

                     Class Level Information

           Class              Levels    Values

           treatment               7    F12 F3 F6 O S12 S3 S6


                 Number of Observations Read        32
                 Number of Observations Used        32
```

Dependent Variable: scabp

```
                                    Sum of
          Source              DF        Squares    Mean Square    F Value    Pr > F
          Model                6     972.343750     162.057292       3.61    0.0103
          Error               25    1122.875000      44.915000
          Corrected Total     31    2095.218750

                 R-Square     Coeff Var     Root MSE    scabp Mean
                 0.464077      42.80633     6.701865      15.65625
```

SAS output

```
           treatment      scabp LSMEAN      95% Confidence Limits

           F12              5.750000       -1.151375    12.651375
           F3               9.500000        2.598625    16.401375
           F6              15.500000        8.598625    22.401375
           0               22.625000       17.744991    27.505009
           S12             14.250000        7.348625    21.151375
           S3              16.750000        9.848625    23.651375
           S6              18.250000       11.348625    25.151375


                          The GLM Procedure

Dependent Variable: scabp

                                Standard
 Parameter                Estimate      Error   t Value  Pr > |t|     95% Confidence Limits

 fall - spring          -6.1666667   2.7360251    -2.25    0.0332   -11.8016158    -0.5317175
 linear trend          -68.1562500  21.0269359    -3.24    0.0034  -111.4620350   -24.8504650
```
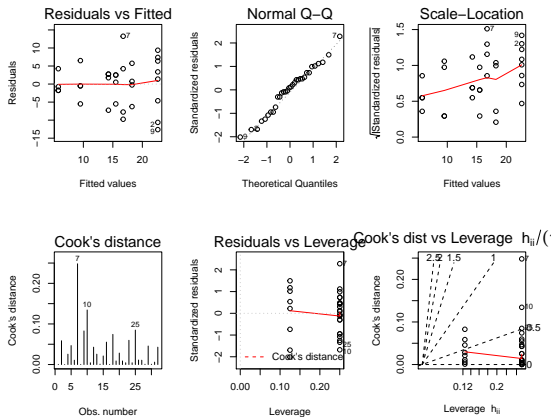
```r
library(multcomp)
K = rbind("fall - spring" = c(1,1,1,0,-1,-1,-1)/3,
          "linear trend"  = c(27,3,-9,-42,-9,3,27)/8)
m = lm(inf~trt,d)
anova(m)


Analysis of Variance Table

Response: inf
          Df Sum Sq Mean Sq F value Pr(>F)
trt        6    972   162.1    3.61   0.01 *
Residuals 25   1123    44.9
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```r
g = glht(lm(inf~trt-1,d), linfct=K) # notice the -1 in the model
summary(g, test=adjusted(type="none")) # unadjusted pvalues
```

```
 Simultaneous Tests for General Linear Hypotheses

Fit: lm(formula = inf ~ trt - 1, data = d)

Linear Hypotheses:
                 Estimate Std. Error t value Pr(>|t|)
fall - spring == 0   -6.17       2.74   -2.25   0.0332 *
linear trend == 0   -68.16      21.03   -3.24   0.0034 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
(Adjusted p values reported -- none method)
```

```r
confint(g, calpha=univariate_calpha()) # unadjusted confidence intervals
```

```
 Simultaneous Confidence Intervals

Fit: lm(formula = inf ~ trt - 1, data = d)

Quantile = 2.06
95% confidence level


Linear Hypotheses:
                 Estimate     lwr      upr
fall - spring == 0   -6.167  -11.802   -0.532
linear trend == 0   -68.156 -111.462  -24.850
```

Summary

For this particular data analysis

- Significant differences in means between the groups (ANOVA $F_{6,25} = 3.61$ p=0.01)
- Fall more effective (p=0.03, 95% CI (0.53, 11.8))
- Linear trend in sulfur (p=0.003)

- Concerned about spatial correlation among columns
- Consider a transformation of the response
  - CI for F12 (-1.2, 12.7)
  - Non-constant variance (residuals vs predicted, sulfur, application)