

Set S05 - Random Forests

STAT 401 (Engineering) - Iowa State University

April 26, 2017

Regression trees

Consider a regression model that uses a set of indicator variables to group the data, e.g.


$$Y_i \stackrel{ind}{\sim} N(\mu_i, \sigma^2)$$

where

$$\begin{aligned} \mu_i = & \beta_0 && \text{group 1} \\ & + \beta_1 \mathbf{I}(x_{i1} \leq 49) \mathbf{I}(x_{i2} > 0.4) && \text{group 2} \\ & + \beta_2 \mathbf{I}(x_{i1} > 49) \mathbf{I}(x_{i3} \leq 0.7) && \text{group 3} \\ & + \beta_3 \mathbf{I}(x_{i1} > 49) \mathbf{I}(x_{i3} > 0.7) && \text{group 4} \end{aligned}$$

Thus group 1 corresponds to those observations with $x_{i1} \leq 49$ and $x_{i2} \leq 0.4$.

Visualization of a regression tree



Regression trees in R tree

```
library("tree")
data(cpus, package="MASS")
m_tree <- tree(log10(perf) ~ syct+mmin+mmax+cach+chmin+chmax, cpus)
summary(m_tree)
```

Regression tree:

```
tree(formula = log10(perf) ~ syct + mmin + mmax + cach + chmin +
      chmax, data = cpus)
```

Variables actually used in tree construction:

```
[1] "cach" "mmax" "syct" "chmin"
```


Number of terminal nodes: 10

Residual mean deviance: 0.03187 = 6.342 / 199

Distribution of residuals:

| Min. | 1st Qu. | Median | Mean | 3rd Qu. | Max. |
|------------|------------|-----------|-----------|-----------|-----------|
| -0.4945000 | -0.1191000 | 0.0003571 | 0.0000000 | 0.1141000 | 0.4680000 |

```
plot(m_tree); text(m_tree)
```



Regression trees in R rpart

```
library("rpart")
m_rpart <- rpart(log10(perf) ~ syct+mmin+mmax+cach+chmin+chmax, cpus)
summary(m_rpart)
```

Call:

```
rpart(formula = log10(perf) ~ syct + mmin + mmax + cach + chmin +
      chmax, data = cpus)
n= 209
```

| | CP | nsplit | rel error | xerror | xstd |
|---|------------|--------|-----------|-----------|------------|
| 1 | 0.54926971 | 0 | 1.0000000 | 1.0080363 | 0.09735912 |
| 2 | 0.08933901 | 1 | 0.4507303 | 0.4701784 | 0.04776144 |
| 3 | 0.08763324 | 2 | 0.3613913 | 0.4274450 | 0.04457527 |
| 4 | 0.03281589 | 3 | 0.2737580 | 0.3227759 | 0.03101707 |
| 5 | 0.02692205 | 4 | 0.2409421 | 0.3118627 | 0.03024666 |
| 6 | 0.01855609 | 5 | 0.2140201 | 0.2954596 | 0.02917108 |
| 7 | 0.01679918 | 6 | 0.1954640 | 0.2919951 | 0.03094696 |
| 8 | 0.01579084 | 7 | 0.1786648 | 0.2873176 | 0.03034303 |
| 9 | 0.01000000 | 9 | 0.1470831 | 0.2588373 | 0.02846206 |

Variable importance

| cach | mmax | mmin | chmin | syct | chmax |
|------|------|------|-------|------|-------|
| 25 | 20 | 17 | 15 | 14 | 9 |

Node number 1: 209 observations, complexity param=0.5492697

mean=1.753333, MSE=0.2062945


left son=2 (143 obs) right son=3 (66 obs)

Primary splits:

cach < 27 to the left, improve=0.5492697, (0 missing)

mmax < 14000 to the left, improve=0.4942141, (0 missing)

```
plot(m_rpart); text(m_rpart)
```



How do these approaches decide on the splits?

From the help file for tree:

A tree is grown by binary recursive partitioning using the response in the specified formula and choosing splits from the terms of the right-hand-side. Numeric variables are divided into $X < a$ and $X > a$; the levels of an unordered factor are divided into two non-empty groups. The split which maximizes the reduction in impurity is chosen, the data set split and the process repeated. Splitting continues until the terminal nodes are too small or too few to be split.


The *impurity* for a regression tree is most likely the estimate of $\hat{\sigma}^2$. Thus, the algorithm searches over all possible splits and finds the one that results in the smallest $\hat{\sigma}^2$. Then the process is repeated for each split.

To determine when to stop, the algorithm has a set of control values. For tree the values are

- mincut: minimum number of observations to include in either child node
- minsize: smallest allowed node size
- mindev: within-node deviance must be at least this times that of the root node for the node to be split

Little tree

```
m_tree <- tree(log10(perf) ~ syct+mmin+mmax+cach+chmin+chmax, cpus,  
  control = list(mincut = 100, mindev = 0.01, minsize = 200, nmax = 90))  
plot(m_tree); text(m_tree)
```



Random forests


Repeat this algorithm B times:

1. Randomly sample data with replacement from training set.
2. Train a tree on these data (randomly evaluating a subset of explanatory variables for each split).
3. Evaluate the tree based on its out of sample performance.

After training, predictions for new data are averaged across all the trees.

Visualizing

Ensemble Model:
example for regression



Random forests in R

```
forest <- randomForest(log10(perf) ~ syct+mmin+mmax+cach+chmin+chmax,
                        data = cpus,
                        importance = TRUE)
forest
```


Call:

```
randomForest(formula = log10(perf) ~ syct + mmin + mmax + cach +          chmin + chmax, data = cpus, importance =
              Type of random forest: regression
              Number of trees: 500
No. of variables tried at each split: 2

              Mean of squared residuals: 0.024559
              % Var explained: 88.1
```

Out of bag error

```
plot(forest)
```



Variable importance

```
importance(forest) %>% round(2)
```

| | %IncMSE | IncNodePurity |
|-------|---------|---------------|
| sycl | 16.94 | 3.79 |
| mmin | 17.84 | 5.43 |
| mmax | 31.60 | 11.08 |
| cach | 32.59 | 11.87 |
| chmin | 19.57 | 6.16 |
| chmax | 20.02 | 2.96 |

Prediction

```
new_cpus = cpus %>%
  sample_n(10)
10^predict(forest, new_cpus, interval="confidence") # Performance with no uncertainty
```

| | | | | | | | | | |
|----------|----------|----------|----------|----------|-----------|----------|-----------|-----------|-----------|
| 101 | 134 | 82 | 116 | 165 | 155 | 86 | 6 | 195 | 153 |
| 21.43632 | 33.77890 | 37.36409 | 71.12837 | 47.50400 | 184.26909 | 30.50168 | 315.67296 | 102.32854 | 318.04787 |

Classification trees

