

Set S05 - Random Forests

STAT 401 (Engineering) - Iowa State University

April 26, 2017

Regression trees

Consider a regression model that uses a set of indicator variables to group the data, e.g.


$$Y_i \stackrel{ind}{\sim} N(\mu_i, \sigma^2)$$

where

$$\begin{aligned} \mu_i = & \beta_0 && \text{group 1} \\ & + \beta_1 \mathbf{I}(x_{i1} \leq 49) \mathbf{I}(x_{i2} > 0.4) && \text{group 2} \\ & + \beta_2 \mathbf{I}(x_{i1} > 49) \mathbf{I}(x_{i3} \leq 0.7) && \text{group 3} \\ & + \beta_3 \mathbf{I}(x_{i1} > 49) \mathbf{I}(x_{i3} > 0.7) && \text{group 4} \end{aligned}$$

Thus group 1 corresponds to those observations with $x_{i1} \leq 49$ and $x_{i2} \leq 0.4$.

Visualization of a regression tree



Regression trees in R tree

```
library("tree")
data(cpus, package="MASS")
m_tree <- tree(log10(perf) ~ syct+mmin+mmax+cach+chmin+chmax, cpus)
summary(m_tree)
```

Regression tree:

```
tree(formula = log10(perf) ~ syct + mmin + mmax + cach + chmin +
      chmax, data = cpus)
```

Variables actually used in tree construction:

```
[1] "cach" "mmax" "syct" "chmin"
```


Number of terminal nodes: 10

Residual mean deviance: 0.03187 = 6.342 / 199

Distribution of residuals:

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-0.4945000	-0.1191000	0.0003571	0.0000000	0.1141000	0.4680000

```
plot(m_tree); text(m_tree)
```



Regression trees in R rpart

```
library("rpart")
m_rpart <- rpart(log10(perf) ~ syct+mmin+mmax+cach+chmin+chmax, cpus)
summary(m_rpart)
```

Call:

```
rpart(formula = log10(perf) ~ syct + mmin + mmax + cach + chmin +
      chmax, data = cpus)
n= 209
```

	CP	nsplit	rel error	xerror	xstd
1	0.54926971	0	1.0000000	1.0080363	0.09735912
2	0.08933901	1	0.4507303	0.4701784	0.04776144
3	0.08763324	2	0.3613913	0.4274450	0.04457527
4	0.03281589	3	0.2737580	0.3227759	0.03101707
5	0.02692205	4	0.2409421	0.3118627	0.03024666
6	0.01855609	5	0.2140201	0.2954596	0.02917108
7	0.01679918	6	0.1954640	0.2919951	0.03094696
8	0.01579084	7	0.1786648	0.2873176	0.03034303
9	0.01000000	9	0.1470831	0.2588373	0.02846206

Variable importance

cach	mmax	mmin	chmin	syct	chmax
25	20	17	15	14	9

Node number 1: 209 observations, complexity param=0.5492697

mean=1.753333, MSE=0.2062945


left son=2 (143 obs) right son=3 (66 obs)

Primary splits:

cach < 27 to the left, improve=0.5492697, (0 missing)

mmax < 14000 to the left, improve=0.4942141, (0 missing)

```
plot(m_rpart); text(m_rpart)
```



How do these approaches decide on the splits?

From the help file for tree:

A tree is grown by binary recursive partitioning using the response in the specified formula and choosing splits from the terms of the right-hand-side. Numeric variables are divided into $X < a$ and $X > a$; the levels of an unordered factor are divided into two non-empty groups. The split which maximizes the reduction in impurity is chosen, the data set split and the process repeated. Splitting continues until the terminal nodes are too small or too few to be split.


The *impurity* for a regression tree is most likely the estimate of $\hat{\sigma}^2$. Thus, the algorithm searches over all possible splits and finds the one that results in the smallest $\hat{\sigma}^2$. Then the process is repeated for each split.

To determine when to stop, the algorithm has a set of control values. For tree the values are

- mincut: minimum number of observations to include in either child node
- minsize: smallest allowed node size
- mindev: within-node deviance must be at least this times that of the root node for the node to be split

Little tree

```
m_tree <- tree(log10(perf) ~ syct+mmin+mmax+cach+chmin+chmax, cpus,  
  control = list(mincut = 100, mindev = 0.01, minsize = 200, nmax = 90))  
plot(m_tree); text(m_tree)
```



Random forests


Repeat this algorithm B times:

1. Randomly sample data with replacement from training set.
2. Train a tree on these data (randomly evaluating a subset of explanatory variables for each split).
3. Evaluate the tree based on its out of sample performance.

After training, predictions for new data are averaged across all the trees.

Visualizing

Ensemble Model:
example for regression



Random forests in R

```
forest <- randomForest(log10(perf) ~ syct+mmin+mmax+cach+chmin+chmax,
                        data = cpus,
                        importance = TRUE)
forest
```


Call:

```
randomForest(formula = log10(perf) ~ syct + mmin + mmax + cach +      chmin + chmax, data = cpus, importance =
              Type of random forest: regression
              Number of trees: 500
No. of variables tried at each split: 2

              Mean of squared residuals: 0.02455569
              % Var explained: 88.1
```

Out of bag error

```
plot(forest)
```



Variable importance

```
importance(forest) %>% round(2)
```

	%IncMSE	IncNodePurity
sycl	16.93	3.80
mmin	17.83	5.43
mmax	31.70	11.08
cach	32.68	11.87
chmin	19.58	6.16
chmax	20.13	2.96

Prediction

```
new_cpus = cpus %>%
  sample_n(10)

new_cpus %>%
  bind_cols(data.frame(pred_perf = 10^predict(forest,
                                              new_cpus))) %>%
  arrange(pred_perf)
```

	name	syst	mmin	mmax	cach	chmin	chmax	perf	estperf	pred_perf
1	IBM 370/148	203	1000	2000	0	1	5	24	21	21.43632
2	HONEYWELL DPS 7/55	140	2000	4000	0	3	6	29	28	30.50168
3	MAGNUSON M80/32	100	1000	8000	24	3	6	32	46	33.76231
4	HONEYWELL DPS 6/92	300	1000	4000	8	3	64	38	30	37.35721
5	NCR V8595 II	56	4000	16000	0	1	8	46	78	47.50400
6	IBM 4341-12	185	2000	16000	16	1	6	76	76	71.12837
7	SPERRY 1100/81	50	2000	32000	24	6	26	114	182	102.32854
8	NAS AS/9000 N	48	4000	24000	32	8	24	214	151	183.99017
9	AMDAHL 470V/8	26	8000	32000	64	8	32	318	290	315.67296
10	NAS AS/8060	35	8000	32000	64	8	24	370	270	318.04787

Classification trees

