Parameter estimation

Dr. Jarad Niemi

STAT 544 - Iowa State University

January 24, 2019

Outline

- Parameter estimation
  - Beta-binomial example
  - Point estimation
  - Interval estimation
  - Simulation from the posterior
- Priors
  - Subjective
  - Conjugate
  - Default
  - Improper

Parameter estimation

For point or interval estimation of a parameter θ in a model M based on data y, Bayesian inference is based off

$$p(\theta|y) = \frac{p(y|\theta)p(\theta)}{p(y)} = \frac{p(y|\theta)p(\theta)}{\int p(y|\theta)p(\theta)d\theta} \propto p(y|\theta)p(\theta)$$

where

- $p(\theta)$ is the prior distribution for the parameter,
- $p(\theta|y)$ is the posterior distribution for the parameter,
- $p(y|\theta)$ is the statistical model (or likelihood), and
- $p(y)$ is the prior predictive distribution (or marginal likelihood).

Obtaining the posterior

The hard way:
1. Derive $p(y)$.
2. Derive $p(\theta|y) = p(y|\theta)p(\theta)/p(y)$.

The easy way:
1. Derive $f(\theta) = p(y|\theta)p(\theta)$.
2. Recognize $f(\theta)$ as the kernel of some distribution.

Definition

The kernel of a probability density (mass) function is the form of the pdf (pmf) with any terms not involving the random variable omitted.

For example, $\theta^{a-1}(1-\theta)^{b-1}$ is the kernel of a beta distribution.

Derive the posterior - the hard way

Suppose $Y \sim Bin(n, \theta)$ and $\theta \sim Be(a, b)$, then

$$
\begin{aligned}
p(y) &= \int p(y|\theta)p(\theta)d\theta \\
&= \int \binom{n}{y}\theta^y(1-\theta)^{n-y}\frac{\theta^{a-1}(1-\theta)^{b-1}}{\mathsf{Beta}(a,b)}d\theta \\
&= \binom{n}{y}\frac{1}{\mathsf{Beta}(a,b)}\int \theta^{a+y-1}(1-\theta)^{b+n-y-1}d\theta \\
&= \binom{n}{y}\frac{\mathsf{Beta}(a+y,b+n-y)}{\mathsf{Beta}(a,b)}
\end{aligned}
$$

which is known as the Beta-binomial distribution.

$$
\begin{aligned}
p(\theta|y) &= p(y|\theta)p(\theta)/p(y) \\
&= \binom{n}{y}\theta^y(1-\theta)^{n-y}\frac{\theta^{a-1}(1-\theta)^{b-1}}{\mathsf{Beta}(a,b)}\bigg/\binom{n}{y}\frac{\mathsf{Beta}(a+y,b+n-y)}{\mathsf{Beta}(a,b)} \\
&= \frac{\theta^{a+y-1}(1-\theta)^{b+n-y-1}}{\mathsf{Beta}(a+y,b+n-y)}
\end{aligned}
$$

Thus $\theta|y \sim Be(a + y, b + n - y)$.

Derive the posterior - the easy way

Suppose $Y \sim Bin(n, \theta)$ and $\theta \sim Be(a, b)$, then

$$
\begin{aligned}
p(\theta|y) &\propto p(y|\theta)p(\theta) \\
&\propto \theta^y (1-\theta)^{n-y} \theta^{a-1}(1-\theta)^{b-1} \\
&= \theta^{a+y-1}(1-\theta)^{b+n-y-1}
\end{aligned}
$$

Thus $\theta|y \sim Be(a+y, b+n-y)$.

Interpretation of prior parameters

When constructing the $Be(a, b)$ prior with the binomial likelihood which results in the posterior

$$\theta|y \sim Be(a + y, b + n - y),$$

we can interpret the prior parameters in the following way:

- a: prior successes
- b: prior failures
- $a + b$: prior sample size
- $a/(a + b)$: prior mean

These interpretations may aid in construction of this prior for a given application.

Posterior mean is a weighted average of prior mean and the MLE

The posterior is $\theta|y \sim Be(a+y, b+n-y)$. The posterior mean is

$$
\begin{aligned}
E[\theta|y] \quad &= \frac{a+y}{a+b+n} \\
&= \frac{a}{a+b+n} + \frac{y}{a+b+n} \\
&= \frac{a+b}{a+b+n}\left(\frac{a}{a+b}\right) + \frac{n}{a+b+n}\left(\frac{y}{n}\right)
\end{aligned}
$$

Thus, the posterior mean is a weighted average of the prior mean $a/(a+b)$ and the MLE y/n with weights equal to the prior sample size $(a+b)$ and the data sample size (n).

Example data

Assume $Y \sim Bin(n, \theta)$ and $\theta \sim Be(1,1)$ (which is equivalent to
Unif(0,1)). If we observe three successes ($y = 3$) out of ten attempts
($n = 10$). Then our posterior is $\theta|y \sim Be(1+3, 1+10-3) \stackrel{d}{=} Be(4,8)$.
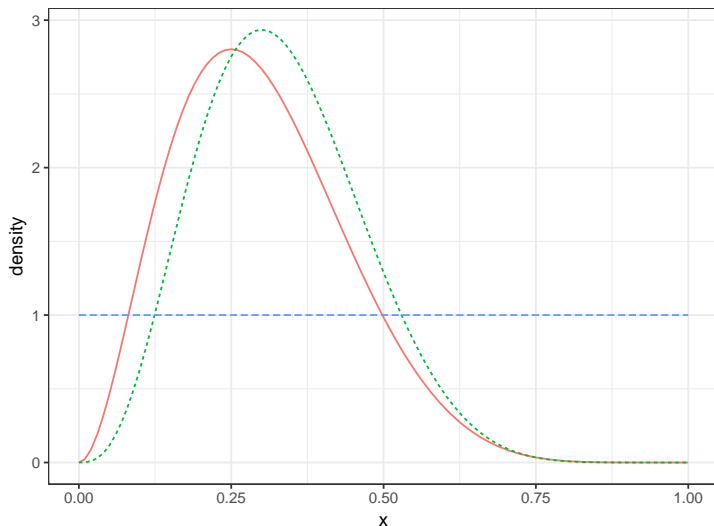The posterior mean is

$$E[\theta|y] = \frac{2}{12} \times \frac{1}{2} + \frac{10}{12} \times \frac{3}{10} = \frac{4}{12}.$$

Remark Note that a $Be(1,1)$ is equivalent to $p(\theta) = \mathrm{I}(0 < \theta < 1)$, i.e.

$$p(\theta|y) \propto p(y|\theta)p(\theta) = p(y|\theta)\mathrm{I}(0 < \theta < 1)$$

so it may seem that a reasonable approach to a default prior is to replace
$p(\theta)$ by a 1 (times the parameter constraint). We will see later that this
depends on the parameterization.

Posterior distribution

Point and interval estimation

Nothing inherently Bayesian about obtaining point and interval estimates.

Point estimation requires specifying a loss (or utility) function.

A $100(1-a)\%$ credible interval is any interval in the posterior that contains the parameter with probability $(1-a)$.

Point estimation

Define a loss (or utility) function $L\left(\theta, \hat{\theta}\right) = -U\left(\theta, \hat{\theta}\right)$ where

- θ is the parameter of interest
- $\hat{\theta} = \hat{\theta}(y)$ is the estimator of θ.

Find the estimator that minimizes the expected loss:

$$\hat{\theta}_{Bayes} = \mathsf{argmin}_{\hat{\theta}} \, E\left[L\left(\theta, \hat{\theta}\right) \middle| y \right]$$

or maximizes expected utility.

Common estimators:

- Mean: $\hat{\theta}_{Bayes} = E[\theta|y]$ minimizes $L\left(\theta, \hat{\theta}\right) = \left(\theta - \hat{\theta}\right)^2$
- Median: $\int_{\hat{\theta}_{Bayes}}^{\infty} p(\theta|y)d\theta = \frac{1}{2}$ minimizes $L\left(\theta, \hat{\theta}\right) = \left|\theta - \hat{\theta}\right|$
- Mode: $\hat{\theta}_{Bayes} = \mathsf{argmax}_{\theta} \, p(\theta|y)$ is obtained by minimizing $L\left(\theta, \hat{\theta}\right) = -\mathrm{I}\left(|\theta - \hat{\theta}| < \epsilon\right)$ as $\epsilon \to 0$, also called maximum a posterior (MAP) estimator.

Mean minimizes squared-error loss

Theorem

The mean minimizes expected squared-error loss.

Proof.

Suppose $L\left(\theta, \hat{\theta}\right) = \left(\theta - \hat{\theta}\right)^2 = \theta^2 - 2\theta\hat{\theta} + \hat{\theta}^2$, then

$$E\left[L\left(\theta, \hat{\theta}\right)\Big| y\right] \qquad = E\left[\theta^2|y\right] - 2\hat{\theta}E[\theta|y] + \hat{\theta}^2$$

$$\frac{d}{d\hat{\theta}}E\left[L\left(\theta, \hat{\theta}\right)\Big| y\right] \quad = -2E[\theta|y] + 2\hat{\theta} \stackrel{set}{=} 0 \implies \hat{\theta} = E[\theta|y]$$

$$\frac{d^2}{d\hat{\theta}^2}E\left[L\left(\theta, \hat{\theta}\right)\Big| y\right] \quad = 2$$

So $\hat{\theta} = E[\theta|y]$ minimizes expected squared-error loss. $\qquad\qquad\square$
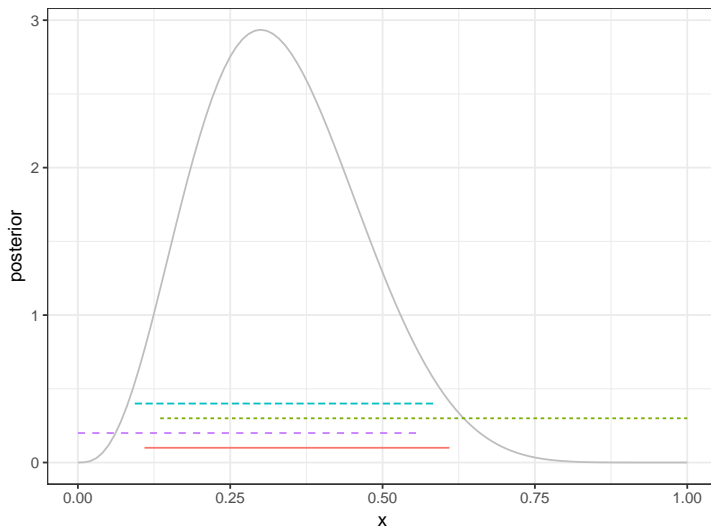
Point estimation
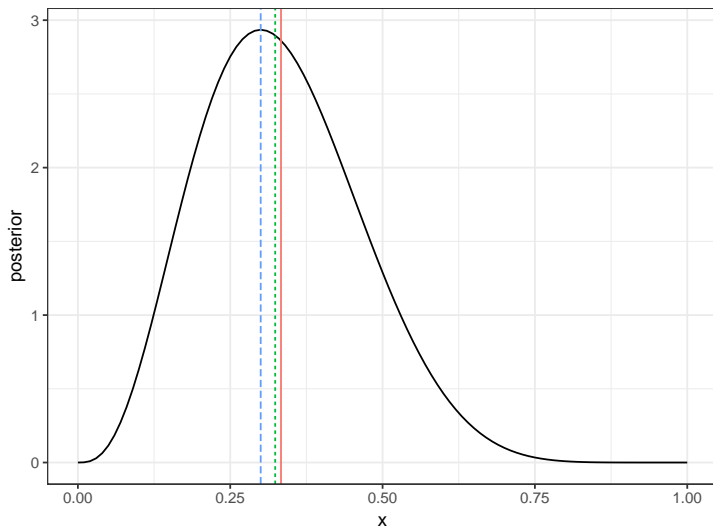
Interval estimation

Definition

A $100(1-a)\%$ credible interval is any interval (L,U) such that

$$1 - a = \int_L^U p(\theta|y)d\theta.$$

Some typical intervals are

- Equal-tailed: $a/2 = \int_{-\infty}^L p(\theta|y)d\theta = \int_U^\infty p(\theta|y)d\theta$
- One-sided: either $L = -\infty$ or $U = \infty$
- Highest posterior density (HPD): $p(L|y) = p(U|y)$ for a uni-modal posterior which is also the shortest interval

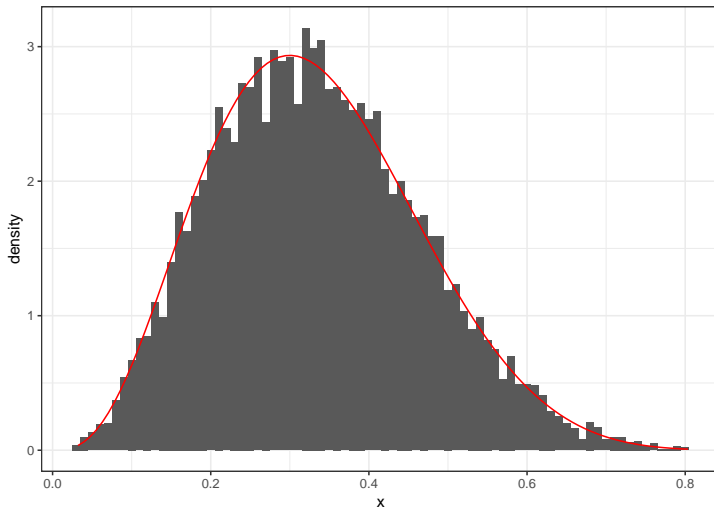# Interval estimation

Simulation from the posterior

An estimate of the full posterior can be obtained via simulation, i.e.

Estimates via simulation

We can also obtain point and interval estimates using these simulations

```
c(mean = mean(sim$x), median = median(sim$x))


     mean    median
0.3351243 0.3270852


quantile(sim$x, c(.025,.975)) # Equal-tail


    2.5%      97.5%
0.1095259 0.6112486


quantile(sim$x, .05) # Upper


       5%
0.1347931


quantile(sim$x, .95) # Lower


      95%
0.5653678
```

Guess the probability

- A coin spins heads.
- New England Patriots win 2019 Super Bowl.
- The first base pair on my genome is A.

What are priors?

Definition

A prior probability distribution, often called simply the prior, of an uncertain quantity θ is the probability distribution that would express one's uncertainty about θ before the "data" is taken into account.

`http://en.wikipedia.org/wiki/Prior_distribution`

Priors

Definition

A prior $p(\theta)$ is conjugate if for $p(\theta) \in \mathcal{P}$ and $p(y|\theta) \in \mathcal{F}$, $p(\theta|y) \in \mathcal{P}$
where $\mathcal{F}$ and $\mathcal{P}$ are families of distributions.

For example, the beta distribution $(\mathcal{P})$ is conjugate to the binomial
distribution with unknown probability of success $(\mathcal{F})$ since

$$\theta \sim \text{Be}(a, b) \qquad \text{and} \qquad \theta|y \sim \text{Be}(a + y, b + n - y).$$

Definition

A natural conjugate prior is a conjugate prior that has the same functional
form as the likelihood.

For example, the beta distribution is a natural conjugate prior since

$$p(\theta) \propto \theta^{a-1}(1 - \theta)^{b-1} \qquad \text{and} \qquad L(\theta) \propto \theta^y(1 - \theta)^{n-y}.$$

Discrete priors are conjugate

Theorem
Discrete priors are conjugate.

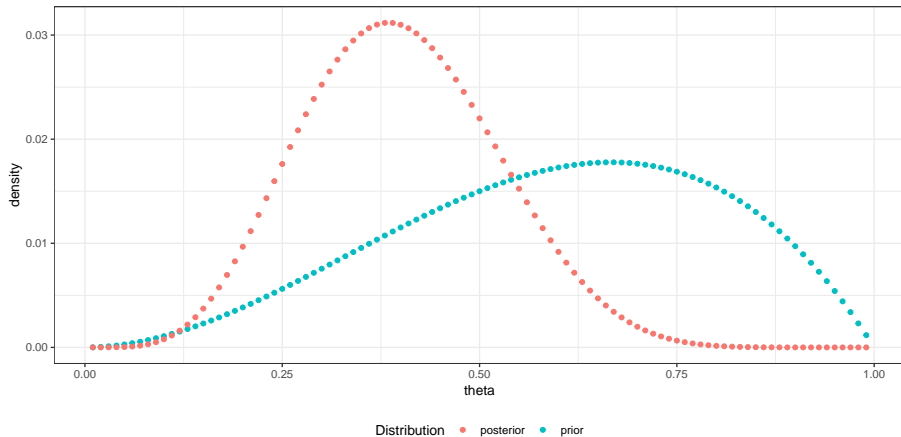## Proof.
Suppose $p(\theta)$ is discrete, i.e.

$$P(\theta = \theta_i) = p_i \qquad \sum_{i=1}^{\mathrm{I}} p_i = 1$$

and $p(y|\theta)$ is the model. Then, $P(\theta = \theta_i|y) = p_i'$ is the posterior with

$$p_i' = \frac{p_i p(y|\theta_i)}{\sum_{j=1}^{\mathrm{I}} p_j p(y|\theta_j)} \propto p_i p(y|\theta_i).$$

$\square$

Discrete prior

Discrete mixtures of conjugate priors are conjugate

Theorem
Discrete mixtures of conjugate priors are conjugate.

Proof.
Let $p_i = P(H_i)$ and $p_i(\theta) = p(\theta|H_i)$,

$$\theta \sim \sum_{i=1}^{I} p_i p_i(\theta) \qquad \sum_{i=1}^{I} p_i = 1,$$

and $p_i(y) = \int p(y|\theta) p_i(\theta) d\theta$, then

$$
\begin{aligned}
p(\theta|y) &= \frac{1}{p(y)} p(y|\theta) p(\theta) = \frac{1}{p(y)} p(y|\theta) \sum_{i=1}^{I} p_i p_i(\theta) \\
&= \frac{1}{p(y)} \sum_{i=1}^{I} p_i p(y|\theta) p_i(\theta) = \frac{1}{p(y)} \sum_{i=1}^{I} p_i p_i(y) p_i(\theta|y) \\
&= \sum_{i=1}^{I} \frac{p_i p_i(y)}{p(y)} p_i(\theta|y) = \sum_{i=1}^{I} \frac{p_i p_i(y)}{\sum_{j=1}^{I} p_j p_j(y)} p_i(\theta|y)
\end{aligned}
$$

Mixtures of conjugate priors are conjugate

Bottom line: if

$$\theta \sim \sum_{i=1}^{I} p_i p_i(\theta) \qquad \sum_{i=1}^{I} p_i = 1$$

and $p_i(y) = \int p(y|\theta)p_i(\theta)d\theta$, then

$$\theta|y \sim \sum_{i=1}^{I} p_i' p_i(\theta|y) \qquad p_i' \propto p_i p_i(y)$$

where $p_i(\theta|y) = p(y|\theta)p_i(\theta)/p_i(y)$.

Mixture of beta distributions

Recall, if $Y \sim Bin(n, \theta)$ and $\theta \sim Be(a, b)$, then the marginal likelihood is

$$
\begin{aligned}
p(y) &= \int p(y|\theta) p(\theta) d\theta \\
&= \int \binom{n}{y} \theta^y (1-\theta)^{n-y} \frac{\theta^{a-1}(1-\theta)^{b-1}}{\mathsf{Beta}(a,b)} \\
&= \binom{n}{y} \frac{1}{\mathsf{Beta}(a,b)} \int \theta^{a+y-1}(1-\theta)^{b+n-y-1} d\theta \\
&= \binom{n}{y} \frac{\mathsf{Beta}(a+y, b+n-y)}{\mathsf{Beta}(a,b)} \quad y = 0, \dots, n
\end{aligned}
$$

which is called the beta-binomial distribution with parameters $a + y$ and $b + n - y$.
If $Y \sim Bin(n, \theta)$ and

$$
\theta \sim p\, \mathsf{Be}(a_1, b_1) + (1-p)\mathsf{Be}(a_2, b_2),
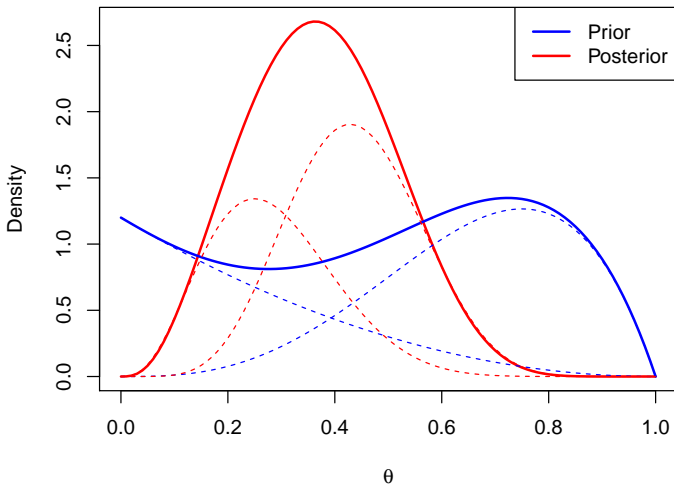$$

then

$$
\theta|y \sim p'\, \mathsf{Be}(a_1 + y, b_1 + n - y) + (1-p')\mathsf{Be}(a_2 + y, b_2 + n - y)
$$

with

$$
p' = \frac{p\, p_1(y)}{p\, p_1(y) + (1-p)p_2(y)} \qquad p_i(y) = \binom{n}{y} \frac{\mathsf{Beta}(a_i + y, b_i + n - y)}{\mathsf{Beta}(a_i, b_i)}
$$

Mixture priors

Binomial, mixture of betas

Default priors

Definition

A default prior is used when a data analyst is unable or unwilling to specify an informative prior distribution.

Default priors

Can we always use $p(\theta) \propto 1$?

Suppose we use $\phi = \log(\theta/[1-\theta])$, the log odds as our parameter, and set $p(\phi) \propto 1$, then the implied prior on θ is

$$
\begin{aligned}
p_\theta(\theta) \propto & \quad 1 \left| \frac{d}{d\theta} \log(\theta/[1-\theta]) \right| \\
= & \frac{1-\theta}{\theta} \left[\frac{1}{1-\theta} + \frac{\theta}{[1-\theta]^2} \right] \\
= & \frac{1-\theta}{\theta} \left[\frac{[1-\theta]+\theta}{[1-\theta]^2} \right] \\
= & \theta^{-1}[1-\theta]^{-1}
\end{aligned}
$$

a Be(0,0), if that were a proper distribution, and is different from setting $p(\theta) \propto 1$ which results in the Be(1,1) prior. Thus, the constant prior is not invariant to the parameterization used.

Fisher information background

Definition

Fisher information, $\mathcal{I}(\theta)$, for a scalar parameter θ is the expectation of the second derivative of the log-likelihood, i.e.

$$\mathcal{I}(\theta) = E\left[\left.\frac{\partial^2}{\partial \theta^2} \log p(y|\theta)\right| \theta\right].$$

Theorem (Casella & Berger (2nd ed) Lemma 7.3.11)

For exponential families,

$$\mathcal{I}(\theta) = -E\left[\left.\left(\frac{\partial}{\partial \theta} \log p(y|\theta)\right)^2\right| \theta\right]$$

If $\theta = (\theta_1, \ldots, \theta_n)$, then the Fisher information is the expectation of the Hessian matrix, which has the ith row and jth that is the partial derivative with respect to θ_i followed by the partial derivative with respect to θ_j, of the log-likelihood.

Jeffreys prior

Definition

Jeffreys prior is a prior that is invariant to parameterization and is obtained via

$$p(\theta) \propto \sqrt{\det \mathcal{I}(\theta)}$$

where $\mathcal{I}(\theta)$ is the Fisher information.

For example, for a binomial distribution $\mathcal{I}(\theta) = \frac{n}{\theta[1-\theta]}$, so

$$p(\theta) \propto \theta^{-1/2}(1 - \theta)^{-1/2} = \theta^{1/2-1}(1 - \theta)^{1/2-1}$$

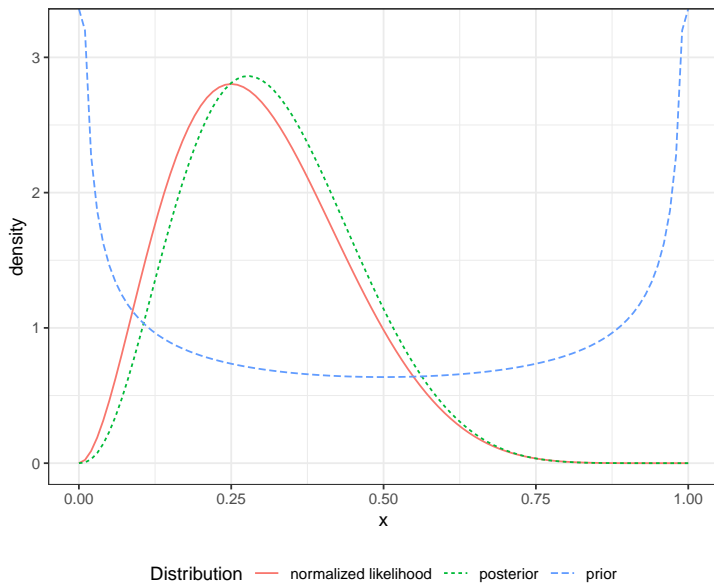a Be(1/2,1/2) distribution.

Fisher information

Theorem

The Fisher information for $Y \sim Bin(n, \theta)$ is $\mathcal{I}(\theta) = \frac{n}{\theta(1-\theta)}$.

Proof.

Since the binomial is an exponential family,

$$
\begin{aligned}
\mathcal{I}(\theta) &= -E_{y|\theta}\left[\frac{\partial^2}{\partial\theta^2}\log p(y|\theta)\right] \\
&= -E_{y|\theta}\left[\frac{\partial^2}{\partial\theta^2}\log\binom{n}{y} + y\log\theta + (n-y)\log(1-\theta)\right] \\
&= -E_{y|\theta}\left[\frac{\partial}{\partial\theta}\frac{y}{\theta} - \frac{n-y}{1-\theta}\right] \\
&= -E_{y|\theta}\left[-\frac{y}{\theta^2} - \frac{n-y}{(1-\theta)^2}\right] \\
&= -\left[-\frac{n\theta}{\theta^2} - \frac{n-n\theta}{(1-\theta)^2}\right] = \frac{n}{\theta} + \frac{n}{(1-\theta)} \\
&= \frac{n}{\theta(1-\theta)}
\end{aligned}
$$

$\square$

Non-conjugate priors

If $Y \sim Bin(n, \theta)$ and $p(\theta) = e^{\theta}/(e-1)$, then

$$p(\theta|y) \propto f(\theta) = \theta^y (1-\theta)^{n-y} e^{\theta}$$

which is not a known distribution.

Options

- Plot $f(\theta)$ (possibly multiplying by a constant).
- Find $i = \int f(\theta)d\theta$, so that $p(\theta|y) = f(\theta)/i$.
- Evaluate $f(\theta)$ on a grid and normalize by the grid spacing.

Plot of $f(\theta)$

Binomial, nonconjugate prior

Numerical integration

Find $i = \int f(\theta)d\theta$, so that $p(\theta|y) = f(\theta)/i$.
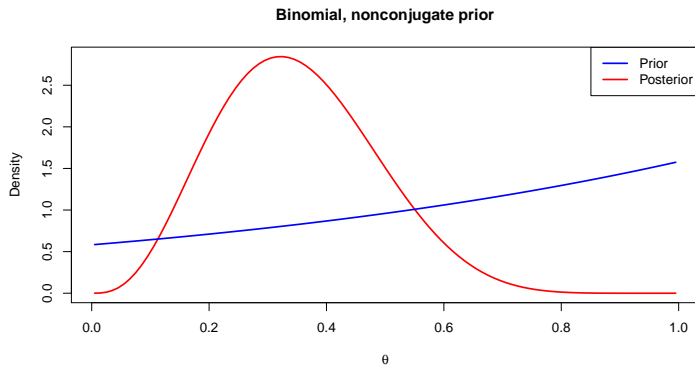
```
(i = integrate(f, 0, 1))

0.001066499 with absolute error < 1.2e-17
```

Nonconjugate prior, numerical integration

Binomial, nonconjugate prior

Nonconjugate prior, evaluated on a grid

Binomial, nonconjugate prior

```
theta[c(which(cumsum(d)*w>0.025)[1]-1, which(cumsum(d)*w>0.975)[1])] # 95\% CI

[1] 0.105 0.625
```

Improper priors

Definition

An unnormalized density, $f(\theta)$, is proper if $\int f(\theta)d\theta = c < \infty$, and otherwise it is improper.

To create a normalized density from a proper unnormalized density, use

$$p(\theta|y) = \frac{f(\theta)}{c}$$

to see that $p(\theta|y)$ is a proper normalized density note that $c = \int f(\theta)d\theta$ is not a function of θ, then

$$\int p(\theta|y)d\theta = \int \frac{f(\theta)}{\int f(\theta)d\theta}d\theta = \int \frac{f(\theta)}{c}d\theta = \frac{1}{c}\int f(\theta)d\theta = \frac{c}{c} = 1$$

Be(0,0) prior

Recall that $Be(a, b)$ is a proper probability distribution if $a > 0, b > 0$.

Suppose $Y \sim Bin(n, \theta)$ and $p(\theta) \propto \theta^{-1}(1 - \theta)^{-1}$, i.e. the kernel of a $Be(0, 0)$ distribution. This is an improper distribution.

The posterior, $\theta|y \sim Be(y, n - y)$, is proper if $0 < y < n$.