Midterm review

Dr. Jarad Niemi

Iowa State University

March 8, 2016

What we have covered

Chapters

- Probability and inference (Ch 1)
- Single-parameter models (Ch 2)
- Introduction to multiparameter models (Ch 3)
- Asymptotics and connections to non-Bayesian approaches (Ch 4)
- Hierarchical models (Ch 5)
- Model checking (Ch 6)
- Bayesian hypothesis tests (Sec 7.4)
- Decision theory (Sec 9.1)
- Stan

Probability and inference (Ch 1)

- Three steps of Bayesian data analysis (Sec 1.1)
  - Set up a full probability model: $p(y|\theta)$ and $p(\theta)$
  - Condition on observed data: $p(\theta|y)$
  - Evaluate the fit of the model: $p(y^{rep}|y)$
- Bayesian inference via Bayes' rule (Sec 1.3)
  - Parameter posteriors: $p(\theta|y) \propto p(y|\theta)p(\theta)$
  - Predictions: $p(\tilde{y}|y) = \int p(\tilde{y}|\theta)p(\theta|y)d\theta$
  - Model probabilities $p(M|y) \propto p(y|M)p(M)$ where $p(y|M) = \int p(y|\theta, M)p(\theta|M)d\theta$.
- Interpreting Bayesian probabilities (Sec 1.5)
  - Epistemic probability: my belief
  - Frequency probability: long run percentage
- Computation (Sec 1.9)
  - Inference via simulations

Single-parameter models (Ch 2)

General

- Priors
  - Conjugate (Sec 2.4)
  - Default - Jeffreys (Sec 2.8)
  - Weakly informative (Sec 2.9)
- Posteriors
  - Compromise between data and prior (2.2)
  - Point estimation
  - Credible intervals (Sec 2.3)

Specific models

- Binomial (Sec 2.1–2.4)
- Normal, unknown mean (Sec 2.5)
- Normal, unknown variance (Sec 2.6)
- Poisson (Sec 2.6)
- Exponential (Sec 2.6)
- Poisson with exposure (Sec 2.7)

Single-parameter models (Ch 2)

Additional comments:

- Deriving posteriors using the kernel
- Discrete priors are conjugate
- Mixtures of conjugate priors are conjugate
- Point estimation depends on utility function
  - Mean minimizes squared error
  - Median minimizes absolute error
  - Mode minimizes 0-1 error
- Computation, $E[h(\theta)|y]$
  - SLLN
  - CLT

Introduction to multiparameter models (Ch 3)

- Joint posterior

$$p(\theta_1, \ldots, \theta_n|y) \propto p(y|\theta_1, \ldots, \theta_n)p(\theta_1, \ldots, \theta_n)$$

- Marginal posterior

$$p(\theta_1|y) = \int \cdots \int p(\theta_1, \ldots, \theta_n|y)d\theta_2 \cdots d\theta_n$$

- Conditional posteriors

$$p(\theta_2, \ldots, \theta_n|\theta_1, y) \propto p(\theta_1, \ldots, \theta_n|y)$$

- Posterior decomposition, e.g.

$$p(\theta_1, \ldots, \theta_n|y) = p(\theta_1|y) \prod_{i=2}^{n} p(\theta_i|\theta_{1:i-1}, y)$$

  where $1 : i-1 = 1, 2, \ldots, i-1$.

- Conditional independence, e.g.

$$p(\theta_i|\theta_{1:i-1}, y) = p(\theta_i|\theta_{i-1}, y)$$

Normal model

- Normal model with default prior (Sec 3.2)

$$y_i \overset{iid}{\sim} N(\mu, \sigma^2) \quad p(\mu, \sigma^2) \propto 1/\sigma^2$$

results in

$$p(\mu, \sigma^2|y) = N(\bar{y}, \sigma^2/n)\text{Inv-}\chi^2(n-1, s^2)$$

where $s^2 = \frac{1}{n}\sum_{i=1}^{n}(y_i - \bar{y})^2$.

- Normal model with conjugate prior (Sec 3.3)

$$y \overset{iid}{\sim} N(\mu, \sigma^2) \quad \mu|\sigma^2 \sim N(\mu_0, \sigma^2/\kappa_0) \quad \sigma^2 \sim \text{Inv-}\chi^2(\nu_0, \sigma_0^2)$$

results in

$$p(\mu, \sigma^2|y) = N\left(\frac{\kappa_0\mu_0 + n\bar{y}}{\kappa_0 + n}, \frac{\sigma^2}{\kappa_0 + n}\right)\text{Inv-}\chi^2(\nu_0 + n, \sigma_n^2)$$

where $\sigma_n^2 = \left[\nu_0\sigma_0^2 + (n-1)s^2 + \frac{\kappa_0 n}{\kappa_0 + n}(\bar{y} - \mu_0)^2\right]/(\nu_0 + n)$.

Data asymptotics (Ch 4)

Consider a model $y_i \overset{iid}{\sim} p(y|\theta_0)$ for some true value θ_0.

- Posterior convergence:
  If A is a neighborhood of θ_0, then $Pr(\theta \in A|y) \to 1$.
- Point estimation:
$$\hat{\theta}_{Bayes} \to \hat{\theta}_{MLE} \xrightarrow{p} \theta_0$$
- Limiting distribution:
$$\theta|y \xrightarrow{d} N\left(\hat{\theta}, \frac{1}{n}\mathrm{I}(\hat{\theta})^{-1}\right)$$

Asymptotics - What can go wrong?

- Not unique to Bayesian statistics
  - Unidentified parameters
  - Number of parameters increase with sample size
  - Aliasing
  - Unbounded likelihoods
  - Tails of the distribution
  - True sampling distribution is not $p(y|\theta)$
- Unique to Bayesian statistics
  - Improper posterior
  - Prior distributions that exclude the point of convergence
  - Convergence to the edge of the parameter space

Hierarchical models (Ch 5)

- Hierarchical model (Ch 5):

$$p(\theta, \phi|y) \propto p(y|\theta)p(\theta|\phi)p(\phi)$$

- Exchangeability (Sec 5.2)

$$p(y_1, \ldots, y_n) = p(y_{\pi_1}, \ldots, y_{\pi_n})$$

- Hierarchical binomial model (Sec 5.3):

$$y_i \overset{iid}{\sim} Bin(n_i, \theta_i) \quad \theta_i \overset{iid}{\sim} Be(\alpha, \beta)$$

- Hierarchical Poisson (with exposure) model

$$y_i \overset{iid}{\sim} Po(x_i\lambda_i) \quad \lambda_i \overset{iid}{\sim} Ga(\mu\beta, \beta)$$

- Hierarchical normal model (Sec 5.4)

$$y_{ij} \overset{iid}{\sim} N(\mu_j, \sigma_j^2) \quad \mu_j \overset{iid}{\sim} N(\eta, \tau^2) \quad \sigma_j^2 \overset{iid}{\sim} Ga(\alpha, \beta)$$

Hypothesis testing (Section 7.4)

From a Bayesian perspective,

$$\text{Simple: } H_i : \theta = \theta_i \qquad \text{Composite: } H_i : \theta \in (\theta_i, \theta_{i+1}]$$

Treat all simple (or all composite) hypotheses as formal Bayesian parameter estimation. Treat a mix of simple and composite hypotheses as formal Bayesian tests.

Formal Bayesian tests

- require prior probabilities for each hypothesis, $p(H_i)$,
- require priors for parameters in non-point hypotheses, $p(\theta|H_i)$, and
- calculate posterior probabilities $p(H_i|y)$ which depend on
- the marginal likelihood, $p(y|H_i)$.

Model checking (Ch 6)

- Data replications

$$p(y^{rep}|y) = \int p(y^{rep}|\theta)p(\theta|y)d\theta$$

- Graphical posterior predictive checks (Sec 6.4)
- Posterior predictive pvalues (Sec 6.3)

$$p_B = P(T(y^{rep}, \theta) \geq T(y, \theta)|y)$$

for a test statistic $T(y, \theta)$.

Decision theory (Sec 9.1)

In order to make a decision, a utility (or loss) function, i.e. $U(\theta, \delta) = -L(\theta, \delta)$, must be set. Then the optimal Bayesian decision is to maximize expected utility (or minimize expected loss), i.e.

$$\text{argmax}_\delta \int U(\theta, \delta) p(\theta) d\theta$$

where $p(\theta)$ represents your current state of belief, i.e. it could be a prior or a posterior depending on your perspective.

Stan

```
model = "
data {
  int<lower=0> N;
  int<lower=0> n[N];
  int<lower=0> y[N];
  real s;
}
parameters {
  real<lower=0,upper=1> mu;
  real<lower=0> eta;
}
transformed parameters {
  real<lower=0> alpha;
  real<lower=0> beta;
  alpha <- eta *    mu;
  beta  <- eta * (1-mu);
}
model {
  mu  ~ beta(20,30);
  eta ~ lognormal(0,s);
  y   ~ beta_binomial(n,alpha,beta);
}
generated quantities {
  real<lower=0,upper=1> theta[N];
  for (i in 1:N) theta[i] <- beta_rng(alpha+y[i], beta+n[i]-y[i]);
}
"
```