I05 - Confidence intervals

STAT 587 (Engineering)
Iowa State University

September 24, 2020

Exact confidence intervals

The coverage of an interval estimator is the probability the interval will contain the true value of the parameter *when the data are considered to be random*. If an interval estimator has $100(1-a)\%$ coverage, then we call it a $100(1-a)\%$ confidence interval and $1-a$ is the confidence level.

That is, we calculate

$$1 - a = P(L < \theta < U)$$

where L and U are random because they depend on the data. Thus confidence is a statement about the procedure.

Normal model

If $Y_i \stackrel{ind}{\sim} N(\mu, \sigma^2)$ and we assume the default prior $p(\mu, \sigma^2) \propto 1/\sigma^2$, then a $100(1-a)\%$ credible interval for μ is given by

$$\overline{y} \pm t_{n-1,a/2} s/\sqrt{n}.$$

When the data are considered random

$$T_{n-1} = \frac{\overline{Y} - \mu}{S/\sqrt{n}} \sim t_{n-1}(0, 1)$$

thus the probability μ is within our credible interval is

$$\begin{aligned}
&P\left(\overline{Y} - t_{n-1,a/2} S/\sqrt{n} < \mu < \overline{Y} + t_{n-1,a/2} S/\sqrt{n}\right) \\
&\quad = P\left(-t_{n-1,a/2} < \frac{\overline{Y}-\mu}{S/\sqrt{n}} < t_{n-1,a/2}\right) \\
&\quad = P\left(-t_{n-1,a/2} < T_{n-1} < t_{n-1,a/2}\right) \\
&\quad = 1 - a.
\end{aligned}$$

Thus, this $100(1-a)\%$ credible interval is also a $100(1-a)\%$ confidence interval.

Yield data example

Recall the corn yield example from I04 with 9 randomly selected fields in Iowa whose sample average yield is 186 and sample standard deviation is 22. Then a 95% confidence interval for the mean corn yield on Iowa farms is

$$186 \pm 2.31 \times 22/\sqrt{9} = (169, 202).$$

Standard error

The standard error of an estimator is an *estimate* of the standard deviation of the estimator (when the data are considered random).

If $Y \sim Bin(n, \theta)$, then

$$\hat{\theta} = \frac{Y}{n} \qquad \text{has} \qquad SE[\hat{\theta}] = \sqrt{\frac{\hat{\theta}(1-\hat{\theta})}{n}}.$$

If $Y_i \overset{ind}{\sim} N(\mu, \sigma^2)$, then

$$\hat{\mu} = \overline{Y} \qquad \text{has} \qquad SE[\hat{\mu}] = S/\sqrt{n}.$$

Approximate confidence intervals

If an unbiased estimator has an asymptotic normal distribution, then we can construct an approximate $100(1-a)\%$ confidence interval for $E[\hat{\theta}] = \theta$ using

$$\hat{\theta} \pm z_{a/2} SE[\hat{\theta}].$$

where $SE[\hat{\theta}]$ is the standard error of the estimator and $P(Z > z_{a/2}) = a/2$.

This comes from the fact that if $\hat{\theta} \dot{\sim} N(\theta, SE[\hat{\theta}]^2)$, then

$$\begin{aligned}
&P\left(\hat{\theta} - z_{a/2} SE(\hat{\theta}) < \theta < \hat{\theta} + z_{a/2} SE(\hat{\theta})\right) \\
&\quad = P\left(-z_{a/2} < \frac{\hat{\theta}-\theta}{SE(\hat{\theta})} < z_{a/2}\right) \\
&\quad \approx P\left(-z_{a/2} < Z < z_{a/2}\right) \\
&\quad = 1 - a.
\end{aligned}$$

Normal example

If $Y_i \overset{ind}{\sim} N(\mu, \sigma^2)$ and we have the estimator $\hat{\mu} = \overline{Y}$, then

$$E[\hat{\mu}] = \mu \qquad \text{and} \qquad SE[\hat{\mu}] = S/\sqrt{n}$$

Thus an approximate $100(1-a)\%$ confidence interval for $\mu = E[\hat{\mu}]$ is

$$\hat{\mu} \pm z_{a/2} SE[\hat{\mu}] = \overline{Y} \pm z_{a/2} S/\sqrt{n}.$$

Note that this is almost identical to the exact $100(1-a)\%$ confidence interval for μ,

$$\overline{Y} \pm t_{n-1,a/2} S/\sqrt{n}$$

and when n is large $z_{a/2} \approx t_{n-1,a/2}$.

T critical values vs Z critical values

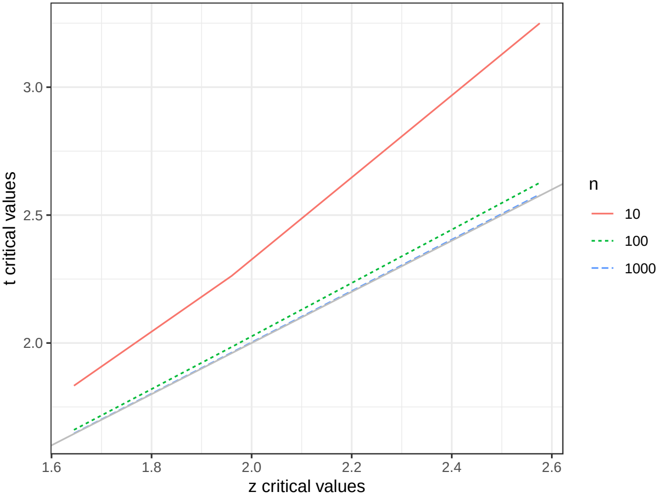

Approximate confidence interval for binomial proportion

If $Y \sim Bin(n, \theta)$, then an approximate $100(1-a)\%$ confidence interval for θ is

$$\hat{\theta} \pm z_{a/2}\sqrt{\frac{\hat{\theta}(1-\hat{\theta})}{n}}.$$

where $\hat{\theta} = Y/n$ since

$$E[\hat{\theta}] = E\left[\frac{Y}{n}\right] = \theta$$

and

$$SE[\hat{\theta}] = \sqrt{\frac{\hat{\theta}(1-\hat{\theta})}{n}}.$$

Gallup poll example

In a Gallup poll dated 2017/02/19, 32.1% of respondents of the 1,500 randomly selected U.S. adults indicated that they were "engaged at work". Thus an approximate 95% confidence interval for the proportion of all U.S. adults is

$$0.321 \pm 1.96 \times \sqrt{\frac{.321(1-.321)}{1500}} = (0.30, 0.34).$$

Confidence interval summary

Model	Parameter	Estimator	Confidence Interval	Type
$Y_i \stackrel{ind}{\sim} N(\mu, \sigma^2)$	μ	$\hat{\mu} = \overline{y}$	$\hat{\mu} \pm t_{n-1,a/2} s/\sqrt{n}$	exact
$Y_i \stackrel{ind}{\sim} N(\mu, \sigma^2)$	μ	$\hat{\mu} = \overline{y}$	$\hat{\mu} \pm z_{a/2} s/\sqrt{n}$	approximate
$Y \sim Bin(n, \theta)$	θ	$\hat{\theta} = y/n$	$\hat{\theta} \pm z_{a/2}\sqrt{\hat{\theta}(1-\hat{\theta})/n}$	approximate
$Y_i \stackrel{ind}{\sim} Ber(\theta)$	θ	$\hat{\theta} = \overline{y}$	$\hat{\theta} \pm z_{a/2}\sqrt{\hat{\theta}(1-\hat{\theta})/n}$	approximate

Bayesian credible intervals generally provide approximate confidence intervals.

Approximate means that the coverage will get closer to the desired probability, i.e. $100(1-a)\%$, as the sample size gets larger.